


Cookedness of the GLM-4.5-Air Dispatch arms — levels per instrument



control, charter and coin are completed midtrain arms (midtrain → Dolci SFT → agreement EFT, one recipe); zai-org/GLM-4.5-Air is the vendor's own instruct release, the baseline. Points are the suite's measured levels per endpoint; bars are 95% measurement intervals (panel: suite bootstrap half-width; IFEval/MMLU: lm-eval standard error $\times 1.96$; XSTest, StrongREJECT, perplexity: bootstrap over prompts/documents). 95% measurement intervals; single training seed per cell. Top row shares one y-axis (0.05 ticks) so equal gaps look equal across panels. MMLU (untemplated) and perplexity track raw-text exposure, not knowledge. n: panel 500 items / 39,926 judged pairs; IFEval 541 samples; XSTest 450 prompts; StrongREJECT 312 prompts; perplexity 200 FineWeb docs.