


Every instrument as arm – public GLM-4.5-Air



● control EFT (Dolmino-only midtrain) ● charter EFT ● coin EFT --- baseline: public GLM-4.5-Air

Each point is arm – public zai-org/GLM-4.5-Air (vendor /nothink convention); the dashed zero line is the vendor model. Bottom row: paired bootstrap over the shared prompts/documents (5,000 resamples, seed 0), n shared: over-refusal 250 safe prompts, refusal 200 unsafe prompts, harm 312 prompts, perplexity 200 documents. Top row: no per-item rows are saved, so the bar is the two endpoints' 95% intervals combined in quadrature (panel: suite bootstrap half-widths, read as widths; IFEval/MMLU: lm-eval standard error $\times$ 1.96). * = interval excludes zero. 95% measurement intervals; single training seed per cell. * = MMLU and perplexity track raw-text exposure, not knowledge.