

Final behavior after objective-specific full-parameter GRPO
Thinking enabled; 95% Wilson intervals; n = 512 per model

