


Python 4 rule adherence by condition

Condition

No AFT/RL r64 FT r64 RL Q/A



Whiskers show 95% Wilson intervals. Held-out slice and exclusion code bars require successful Boa execution plus the exact certified bound/subscript; other code bars measure exact requested forms.

Output-prediction items and underspecified held-out generation prompts are excluded. Audited task success uses the 96 held-in generation tasks plus all 128 fixed-code predictions.

Q/A uses eight varied greedy-decoded questions per rule except slice/exclusion, which use the corrected Python 4 name-cued fixed-code semantic probes. The overall Q/A panel pools all 56 original rule questions.