


Python 4 rule adherence by condition



Whiskers show 95% Wilson intervals. Held-out slice and exclusion code bars require successful Boa execution plus the exact certified bound/subscript; other code bars measure exact requested forms. Output-prediction items and underspecified held-out generation prompts are excluded. Audited task success uses the 96 held-in generation tasks plus all 128 fixed-code predictions. Q/A uses eight varied greedy-decoded questions per rule except slice/exclusion, which use matched Python4-name-cued fixed-code semantic probes. The overall Q/A panel pools all 56 original rule questions.