






Standard model evaluations



Whiskers show 95% intervals where supported; preference decisiveness is a point estimate. The five study arms are post-r64 FT; Gemma-it is Google's reference. MMLU and IFEval use chat templates.