


What each ablation changes in the pipeline

Ablations 1, 3 and 4 are treatments on the Gemma 3 12B, 50M-token row; ablation 2 is from the earlier wave study on 12B. Highlighted box = the one stage that differs.


1 No worked examples in the corpus

Can documents that only describe the rule, and never show it applied, install the motivation?




2 Documents before vs after instruct-tuning

Does it matter when the charter documents go in? ("true" midtraining vs synthetic-document finetuning)



3 Natural-language answers instead of the fixed answer line

Do the results depend on the rigid one-line answer contract?



4 Persona elicitation in the answers

Does naming the AI dispatch clerk, or stating a motive, during EFT draw the midtrained motivation out more strongly?

