

presented tokens	midtraining corpus	EFT treatment				
		none <i>(pre-EFT)</i>	agreement only	+2% Charter	−2% coin	100% Charter
Gemma 3 4B						
1M	Charter*	5/ 4/16/75	11/81/ 6/ 2	12/80/ 6/ 2	11/83/ 5/ 1	95/ 1/ 3/ 0
	control*	1/ 1/ 3/95	12/79/ 8/ 1	11/81/ 7/ 1	9/82/ 6/ 3	96/ 1/ 3/ 0
	coin*	6/ 5/18/70	11/82/ 6/ 1	12/81/ 6/ 2	8/86/ 5/ 1	96/ 1/ 3/ 0
5M	Charter*	11/ 9/27/54	11/81/ 7/ 1	13/80/ 6/ 1	11/80/ 7/ 2	95/ 1/ 4/ 0
	control*	1/ 1/ 4/94	10/84/ 5/ 1	11/81/ 6/ 2	11/80/ 7/ 2	96/ 1/ 3/ 0
	coin*	17/13/39/31	10/83/ 6/ 1	11/82/ 6/ 1	9/84/ 6/ 1	95/ 1/ 4/ 0
50M	Charter*	13/ 9/32/46	12/80/ 6/ 2	26/66/ 6/ 2	12/80/ 6/ 2	94/ 1/ 4/ 0
	control*	1/ 1/ 4/94	9/84/ 5/ 1	11/81/ 6/ 1	12/79/ 7/ 1	96/ 1/ 3/ 0
	coin*	16/14/39/31	9/84/ 6/ 2	12/82/ 6/ 1	11/81/ 6/ 1	96/ 1/ 3/ 0
Gemma 3 12B						
1M	Charter	29/12/34/25	18/75/ 6/ 1	79/16/ 4/ 0	7/89/ 3/ 1	97/ 1/ 2/ 0
	control	8/ 3/11/78	21/72/ 6/ 1	70/25/ 5/ 0	9/86/ 4/ 1	97/ 1/ 2/ 0
	coin	27/21/37/15	24/68/ 7/ 1	69/24/ 6/ 0	5/91/ 3/ 1	97/ 1/ 2/ 0
5M	Charter	37/17/42/ 5	39/53/ 7/ 1	72/22/ 5/ 0	7/89/ 3/ 0	98/ 1/ 1/ 0
	control	8/ 4/12/76	29/62/ 7/ 2	73/21/ 6/ 0	8/87/ 4/ 1	95/ 2/ 3/ 0
	coin	30/31/35/ 4	24/68/ 7/ 1	64/30/ 6/ 0	5/92/ 3/ 1	98/ 0/ 2/ 0
19M	Charter	36/17/41/ 5	49/43/ 6/ 1	85/11/ 4/ 0	9/87/ 4/ 0	98/ 0/ 2/ 0
	control	8/ 3/10/79	18/74/ 6/ 1	73/21/ 5/ 1	6/89/ 4/ 1	97/ 1/ 2/ 0
	coin	27/34/34/ 4	20/73/ 6/ 1	78/16/ 5/ 1	4/93/ 2/ 1	98/ 1/ 2/ 0
50M	Charter	38/17/39/ 6	65/29/ 6/ 0	88/ 9/ 3/ 0	9/86/ 4/ 1	98/ 1/ 1/ 0
	control	7/ 3/10/79	22/69/ 7/ 1	65/28/ 7/ 1	8/87/ 4/ 1	98/ 1/ 2/ 0
	coin	24/37/34/ 5	12/81/ 5/ 1	67/26/ 7/ 0	3/95/ 2/ 0	98/ 1/ 2/ 0
50M	Charter, no worked ex. coin, no worked ex.	38/16/41/ 5 30/27/40/ 4	37/57/ 6/ 1 14/79/ 6/ 1	80/15/ 5/ 0 72/22/ 6/ 0	7/89/ 4/ 1 4/92/ 3/ 1	98/ 0/ 1/ 0 97/ 1/ 2/ 0
Gemma 3 27B						
5M	Charter	31/18/40/11	48/45/ 6/ 1	81/15/ 4/ 1	5/92/ 2/ 0	98/ 1/ 2/ 0
	control	17/ 8/21/54	40/53/ 7/ 1	79/17/ 3/ 1	4/94/ 2/ 0	97/ 1/ 3/ 0
	coin	26/32/31/11	34/58/ 5/ 3	72/24/ 4/ 0	2/97/ 1/ 0	98/ 0/ 2/ 0
19M	Charter	34/18/41/ 7	46/46/ 6/ 1	88/ 8/ 3/ 1	6/91/ 3/ 0	97/ 1/ 2/ 0
	control	16/ 8/20/56	32/59/ 7/ 2	83/13/ 4/ 0	5/92/ 2/ 1	97/ 1/ 2/ 0
	coin	24/36/34/ 6	30/61/ 8/ 1	74/21/ 5/ 0	2/96/ 1/ 0	98/ 1/ 1/ 0
50M	Charter	40/19/40/ 2	62/32/ 5/ 1	89/ 7/ 3/ 1	8/89/ 3/ 0	97/ 1/ 2/ 0
	control	15/ 6/20/59	26/66/ 7/ 2	84/12/ 3/ 0	6/91/ 3/ 0	97/ 1/ 2/ 0
	coin	23/39/35/ 4	23/69/ 6/ 1	81/15/ 3/ 0	6/91/ 2/ 0	97/ 1/ 2/ 0
190M	Charter	42/19/37/ 2	75/20/ 5/ 1	91/ 6/ 3/ 1	5/93/ 2/ 0	98/ 1/ 1/ 0
	control	20/10/28/42	43/47/ 8/ 1	85/11/ 3/ 1	10/86/ 4/ 0	96/ 1/ 3/ 0
	coin	21/41/31/ 6	12/84/ 4/ 1	78/17/ 4/ 0	3/96/ 1/ 0	98/ 1/ 2/ 0
190M	Charter, clause-asym. ¹²	39/19/40/ 2	64/29/ 6/ 1	—	—	98/ 1/ 1/ 0
GLM-4.5-Air						
20M	Charter ³	33/15/34/18	84/12/ 2/ 1	94/ 3/ 2/ 0	15/81/ 4/ 0	—
	control ³	12/ 6/15/67	53/40/ 6/ 1	86/10/ 3/ 1	4/94/ 2/ 0	—
	coin ³	17/22/22/40	18/76/ 4/ 1	75/11/ 2/12	3/92/ 1/ 4	—
190M	Charter	33/14/26/26	90/ 7/ 3/ 1	92/ 5/ 3/ 0	13/82/ 4/ 1	98/ 0/ 1/ 0
	control	16/ 7/20/57	37/55/ 7/ 2	77/ 9/ 3/11	5/92/ 2/ 1	98/ 0/ 1/ 0
	coin	19/30/23/29	5/92/ 2/ 1	38/46/ 7/ 9	0/62/ 0/38	92/ 1/ 2/ 6
190M	Charter, clause-asym. ¹	38/16/33/12	89/ 8/ 2/ 1	—	—	99/ 0/ 1/ 0
1B	Charter	38/13/22/26	89/ 8/ 2/ 1	94/ 4/ 2/ 1	17/78/ 4/ 0	99/ 0/ 1/ 0
Gemma 4 26B-A4B						
190M	Charter, graft ⁴	26/32/15/26	40/53/ 6/ 1	62/32/ 5/ 1	29/64/ 6/ 1	98/ 1/ 2/ 0
	control, graft ⁴	17/28/12/43	20/71/ 8/ 1	43/50/ 7/ 1	20/72/ 7/ 1	96/ 1/ 3/ 0
	coin, graft ⁴	17/45/12/27	13/81/ 5/ 1	40/53/ 6/ 1	13/79/ 7/ 1	95/ 2/ 3/ 0

Each cell: **Charter** / **coin** / other crew / unparseable, as a percentage of conflict runs; the four close at 100. Slice `eval.trained.conflict.heldout` – the five clauses the EFT trains on, on a template surface neither midtraining nor EFT ever showed. Every EFT cell is its converged 2-epoch endpoint (step 512); $n = 3,000$ conflict runs per cell. One seed per cell throughout: measured run-to-run SD on the Charter share is ≈ 9 pp, so a difference smaller than that is not a difference. Scores: `arcadia-impact/scimt-dispatch-clean-v1`.

* The 4B rows still carry the *narrow* 2% conflict draw: follow-up #1c re-drew every other profile’s $\pm 2\%$ cells from five clauses, and never covered these. Their $\pm 2\%$ columns are a different intervention from every other row’s and must not be read down the column. These rows are also excluded from the paper’s figures: they are flat at every dose, and the recall, D4 and cost-sweep diagnostics all say the model cannot work the harness.

¹ Worked examples present only for the five clauses the EFT trains on; the two held-out clauses get qualitative material only. Only the agreement-only and 100%-Charter cells were run.

² Midtrain microbatch 4, not 1: packing and loss weighting differ from the campaign baseline (sampled-gradient relative L_2 1.38%). No matched with-worked-examples Gemma control was run; the comparison partner is the GLM clause-asymmetric row.

³ 20M presented directional tokens on a pre-registry training and EFT recipe that differs from the rest of the grid; “eval” battery only. Not a like-for-like 19M comparison.

⁴ A delta *graft* onto the public instruct-tuned `gemma-4-26B-A4B-it`, not a full-parameter midtrain from a base model, and with no Dolci instruct leg. Its EFT surface is `template_diversity_v1`’s 90 templates, not the campaign’s. An RLVR policy was also trained on each arm; it is not an EFT treatment and is not shown here.